

Uni IT Security Notes

Felix Pojtinger

June 18, 2021

Contents

Uni IT Security Notes	2
Basics	2
Security Mindset	2
Security Objectives	3
Attacks, Threats and Vulnerabilities	3
Threat Identification	3
Security Frameworks	4
Network Specific Threat Examples	4
STRIDE: Attacks on a Multi-User System	4
Security policies	4
Malware	4
Networking	4
TCP Overview	4
TCP Connection Establishment	5
IP Security Issues	5
TCP Security Issues	6
Port Scanning	6
TCP Protection Mechanisms	6
Session Hijacking	7
RST Attacks (In-Connection DoS)	7
Blind IP Spoofing	7
Perimeter Defense in Practice	7
Architecture Recommendations	7
Application in Networking	8
Stateless Packet Filter	8
Stateful Packet Filters	8
Stateful Firewalls	9
Application Layer Proxies	9
Application Level Gateways	9
Demilitarized Zone (DMZ)	10
Web Application Firewalls (WAFs)	10
Intrusion Detection Systems (IDS)	10

Symmetric Encryption	11
Symmetric Encryption Overview	11
Kerckhoffs' Principle	11
Strong Algorithms	11
Crypto Attack Classes	11
Perfect Security	12
One-Time-Pad	12
Stream Cyphers	12
Cryptographically Secure Pseudo-Random Number Generators (CSPRNG)	12
Design Principles for Block Cyphers	12
Feistel Networks	13
DES (Tripple DES)	13
AES Key Features	13
Modes of Operation for Block Cyphers	13
Cypher Block Chaining (CBC)	13
Counter Mode (CTR)	14
Padding	14
Key Length Considerations	14
Message Authentication	15
Message Authentication Codes (MACs)	15
General Scenario	15
Scenario with Modified Message	15
MAC Computation	16
Hash Function Requirements	16
Asymmetric Encryption	16
Public Key Cryptography	16
RSA Key Generation	16
RSA Encryption	17
RSA Decryption	17
RSA Security	17
Hybrid Method	17
Discrete Logarithms	18
One-Way Functions	18
Diffie-Hellman Key Exchange Protocol	18
Diffie-Hellman Key Exchange Protocol Security	18
Digital Signatures	19
RSA Signatures	19

Uni IT Security Notes

Basics

Security Mindset

- Focus on weaknesses, not on features

- Don't rely on the "good case"
- Anticipate what an attacker could do to a system
- Weight security against user experience and privacy

Security Objectives

- **Confidentiality/conf**
 - Nobody but the legitimate receiver can read a message
 - Third party cannot gain access to communication patterns
- **Integrity/int**: The contents of communication can't be changed
- **Authenticity/authN**
 - **Entity Authentication**: Communication partners can prove their respective identity to one another
 - **Message Authentication**: It can be verified that a message is authentic (unaltered and sent by the correct entity)
- **Authorization/authZ**
 - Service or information is only available to those who have correct access rights
 - Depends on authentication being set up
- **Non-Repudiation/nRep**: A sender cannot deny having sent a message or used a service
- **Availability/avail**: Service is available with sufficient performance
- **Access Control/ac**: Access to services and information is controlled
- **Privacy/priv**
 - Restricted access to identity-related data
 - Anonymity
 - Pseudonymity

Attacks, Threats and Vulnerabilities

- **Attacker**: A person who has the skill and motivation to carry out an attack: The steps needed to carry out an attack
- **Vulnerability**: Some characteristics of the target that can result in a security breach
- **Threat**: Combination of an attacker, an attack vector and a vulnerability
- **Attack**: A threat that has been realized and has caused a security breach

Threat Identification

- Define **system boundaries**: What is part of your system, what is not?
- Define **security objectives**: What is important for your system to be secure?
- **List all threats** you can think of: Brainstorming and discussion with experts
- Use **conventions**:
 - Similar threat models
 - Requirement specifications

- How to break or circumvent the specifications
- Note security assumptions of the system
- Be careful with perimeter security: What if perimeter has been breached?
- Note *possible*, but not yet exploitable vulnerabilities

Security Frameworks

Network Specific Threat Examples

- Remote Attacks
- Eavesdropping: Sniffing of information
- Altering information
- Spoofing
- DoS
- Session hijacking
- Viruses attacking clients
- Spam
- Phishing
- Data trails/privacy leaks

STRIDE: Attacks on a Multi-User System

- Spoofing of Identity
- Tampering with Information
- Repudiation
- Information Disclosure
- DoS
- Escalation of Privileges

Security policies

- Classification of system states into “allowed” and “forbidden” states
- Secure system: Is only in allowed states
- Breached system: Is in forbidden state

Malware

- Performs unwanted functions
- Often runs without user’s consent
- Telemetry (often hidden in proprietary software behind EULAs)
- Backdoors

Networking

TCP Overview

- Characteristics

- Reliable
- Connection-Oriented
- Full-Duplex
- Layer atop IP
- Connection management: Setup, Release and Abort
- Ordered delivery (package sequence control)
- Repetition of lost packets
- End-to-End ACKs
- Checksum in header
- Identified by a 5-tuple
 - Source IP
 - Destination IP
 - Transport Protocol
 - Source Port
 - Destination Port

TCP Connection Establishment

- Virtual connection between two systems
- 3-Way-Handshake with connection states

An example connection from the client to the server:

<Client>		<Server>
[Closed]		[Closed]
	SEQ=x CTL=SYN =>	
[SYN Sent]	<= SEQ=y CTL=SYN+ACK ACK=x+1	
	SEQ=x+1 CTL=ACK ACK=y+1 =>	[SYN Received]
[Established]		[Established]

IP Security Issues

- IP header doesn't have confidentiality or integrity protection
 - Faking the sender address is easy to do
 - Traffic can be analyzed by sniffing packet headers
- IP payload doesn't have confidentiality or integrity protection
 - Eavesdropping is possible by sniffing packets
- Loose coupling with lower layers:
 - Easy to divert traffic
 - Availability can be easily attacked
 - Confidentiality and integrity can't be guaranteed
- Unprotected error signaling via ICMP: Fake error messages can affect availability
- DNS is insecure; i.e. DNS spoofing

TCP Security Issues

- TCP header doesn't have confidentiality or integrity protection
- Session hijacking
 - When sniffing session details, attacker can impersonate a peer in a TCP connection
 - Attackers can guess session details and attack remotely using spoofed IP addresses
- RST attack: Attackers can reset/abort attacks by injecting packets with the RST flag
- Port scanning
 - Find out open ports
 - Determine software running on port
- SYN flooding
 - Overload system resources by initializing many connections and not pursuing them

Port Scanning

- Objective: **Collect information**
 - Installed services
 - Software versions
 - OS
 - Firewall
- Enumeration based on port
 - Well-known ports (i.e. SSH → 22)
 - Invalid connection requests: Different way of error handling can be used to fingerprint the OS
- Possible scanning methods
 - TCP connect scan
 - Half-open scan
 - SYN-ACK scan
 - ACK scan

TCP Protection Mechanisms

- SYN flood protection
 - Limit rate of SYN packets
 - SYN cookies (RFC 4987)
 - * Limit resources
 - * Half-open connections are not stored in the connection table but instead as a hash in the ISN
 - * Only if the 3rd ACK handshake packet matches the sequence number, the connection is added to the connection table
 - * Server does not need to maintain any state information on half-open connections: Resources can't be exhausted

- Connections are only accepted if the sequence numbers are within a certain range of acceptable values (attackers would have to sniff sequence numbers or guess them)

Session Hijacking

- Attacker takes over existing connection between two peers
- Requirement: Attacker has to sniff or guess sequence numbers of the connection correctly

RST Attacks (In-Connection DoS)

Inject packet with RST flag into ongoing connection: Connection has to be aborted immediately

Blind IP Spoofing

Firewall is configured to only allow one source IP address and destination IP address ($A \rightarrow B$).

To circumvent this restriction:

1. Attacker starts DoS attack on A to prevent A from sending RST packets to B
2. Attacker sends TCP connection setup packet with A's source IP address to B
3. B sends SYN+ACK packet to A, but can't respond due to DoS
4. Attacker sends TCP connection ACK packet to B with ACK matching the initial sequence number chosen by B (which has to be guessed, as B sent the SYN+ACK packet to A, not the attacker)

Only works if B uses a predictable algorithm for its ISN and packet filters aren't in place.

Perimeter Defense in Practice

Architecture Recommendations

- Known from medieval cities, castles etc.
- Definition of system boundary between "inside" and "outside"
- Different threat models for inside and outside
 - **Inside:** Trusted
 - **Outside:** Untrusted
- Objectives
 - Create said boundary
 - Only a defined set of communication relations is allowed
 - Special security checks
 - Limited number of interconnection points
 - Simpler to manage and audit than a completely open architecture

- Problems
 - Requires intelligent selection of system boundaries
 - May require multiple levels of perimeters
 - No system/user in the “trusted inside” can truly be trusted

Application in Networking

- Installing security devices at the network border
- Separation of network areas into inside/outside
- Prevent sensitive information from being sent to the outside (view the system in the inside as the potential, probably unintentional attacker)
- Multiple levels can increase security
- But: Perimeter security is not sufficient on its own!
 - There will probably be additional non-secured paths into the network (i.e. `ssh -R`)
 - Some malicious traffic might look like “normal” traffic and can pass

Stateless Packet Filter

- Access Control List (ACL): Applies set of rules to each incoming packets
- Discards (denies, blocks) or forwards (allows, permits) packets based on ACL
- Typically configured by IP and TCP/UDP header fields
- Stateless inspection: Established connections can only be detected with the ACK control flag
- Can be easy to misconfigure by forgetting essential protocols
 - DNS
 - ICMP
- Advantages
 - Fast/High throughput
 - Simple to realize
 - Software-based, can be added as a package
 - Simple to configure
- Disadvantages
 - Inflexible
 - Many attacks can only be detected using stateful filtering
 - Rules and their priorities can easily get confusing
- Default discard policy
 - Block everything which is not explicitly allowed (allowlist)
 - Issue: The security policy has to be revised for each new protocol or service
 - This rule must come last/have the lowest priority, behind all “allowing” rules

Stateful Packet Filters

- Store connection states

- Can make decisions based on
 - TCP connections
 - UDP replies to previous outgoing packet with same IP:Port relation (“UDP Connection”)
 - Application protocol states
- Similar to application layer gates/proxy firewalls, but less intruding in communication
- Rules can be more specific than in stateless packet filters
- Rules are easier to enforce, i.e. incoming TCP packets don’t have to be allowed in because they have ACK set

Stateful Firewalls

- Tries to fix the problems of stateless inspection
 - To many packets have to be allowed by default (ACK → No SYN-scanning protection)
 - Protocols like FTP or SIP, which dynamically allocate port numbers, can’t be filtered securely
- Create state per TCP or UDP flow
 - Source and Destination IP:Port
 - Protocol
 - Connection state
- A packet which is not associated with a state is dropped immediately
- Packets which belong to a previously established TCP/UDP “connection” are allowed to pass without further checks
- State tables have to be cleaned up periodically to prevent resource starvation

Application Layer Proxies

- Protected host during connection establishment
- Different kinds
 - Application level
 - Circuit level
 - Forward proxy (client-side)
 - Reverse proxy (server-side)

Application Level Gateways

- Conversion between different application layer protocols
- Evaluation up to OSI layer 7
 - Protocol verification
 - Authentication
 - Malware scanning
 - Spam filtering
 - Attack pattern filtering
- Advantage: Security policies can be enforced at application level

- Disadvantage: Computing and memory performance requirements

Demilitarized Zone (DMZ)

- **Outside world:** Global Internet
- **Outside router:** Routes packet to and from bastion host
- **Bastian host:** Proxy server and relay host
- **Inside router:** Routes packets only to and from bastion host
- **Inside (protected):** Intranet

The DMZ creates 2/3 lines of defense by the use of a stub network.

Multi-Level DMZs can create even more secure perimeter defenses:

Global Internet → Access Router and Packet Filter → Public Services Host (offers i.e. public Web services) → Screening Router and Packet filter (prevents IP spoofing) → Mail host (for external mail communication) → Bastion host (i.e. proxy for FTP and Web access) → Intranet

Web Application Firewalls (WAFs)

- Acts on the application layer
- Is a reverse proxy
- Can protect the web server from “evil” client input
 - Cross-Site scripting
 - SQL injection: Filters out JS or SQL commands in client input by removing special symbols (i.e. <, ' etc)
 - Cookie poisoning: Stores the hash values of sent cookies
 - HTML manipulation: Encrypts URL parameters

Intrusion Detection Systems (IDS)

- Security product that is specialized on detecting anomalies during live operation of networks and computers
 - Virus/Botnet activity
 - Suspicious network activity (malware phoning home)
- Basic Approaches
 - **Signature based:** Use attack signatures/known malicious communication activity patterns
 - **Anomaly based:** Significant deviation from previously recorded baseline activity
 - **Rule based:** Define allowed behaviour by app-specific set of legitimate actions
- Actions
 - Send out alarm
 - Logging
 - Blocking of known patterns
- Realization

- Appliance
- Integration in firewall
- Integration into host

Symmetric Encryption

Symmetric Encryption Overview

Alice:

1. Creates message
2. Chooses key
3. Computes ciphertext
4. Send ciphertext to Bob

Eve (Attacker):

1. Copies ciphertext
2. Tries to guess the key

Bob:

1. Receives ciphertext
2. Uses key
3. Computes plaintext
4. Reads message

Kerckhoffs' Principle

- From “La Cryptographie Militaire”
- Most important point: **The security of a crypto system must lie in the non-disclosure of the key but not in the non-disclosure of the algorithm**
- Implementation
 - Keep secret which function you used for encryption
 - But a disclosure of the set of functions should not create a problem

Strong Algorithms

- There is no attack that can break it with less effort than a brute force attack (“complete enumeration”)
- There are so many keys that a complete search of key space is infeasible

Crypto Attack Classes

- **Active** attacks
 - Most relevant for cryptographic protocols
 - Active interference (modification, insertion or deletion of messages)
 - Man in the middle (MITM) can receive messages and modify them on the way to the receiver

- **Passive attacks:** Pure eavesdropping, without interference with communication

Perfect Security

Ciphertext does not give any information you don't already have about the plaintext

One-Time-Pad

- **Vernam Cypher:** Create ciphertext by XOR addition of secret key and plaintext
- **Mauborgne:** Random key, never re-use key ("one time")
- **Shannon:** OTP is unbreakable if key is ...
 - Truly random
 - As large
 - Never reused
 - Kept secret

Stream Cyphers

Encryption like one-time-pad, but using pseudo-random bits instead of true random (using a **Cryptographically Secure Pseudo-Random Number Generator (CSPRNG)**)

Cryptographically Secure Pseudo-Random Number Generators (CSPRNG)

A CSPRNG must ...

- Be unpredictable
- Be computationally infeasible to compute the next outputs

... when the initial state of the CSPRNG is not known

Design Principles for Block Cyphers

Two methods for frustrating a statistical analysis:

- **Confusion:** The ciphertext should depend on the plaintext in such a complicated way that an attacker cannot gain any information from the ciphertext (redundancy should not be visible anymore in the ciphertext)
- **Diffusion:** Each plaintext and key bit should influence as many ciphertext bits as possible
 - Changing one bit in plaintext → Many pseudo-random changes in ciphertext
 - Changing one bit in the key → Many pseudo-random changes in ciphertext

Feistel Networks

- Described by Horst Feistel
- Algorithm
 - Plaintext block B is divided in 2 halves
 - Derive r round key keys from key
 - Feed one half through round function F
 - Then XOR the result with the other half
 - Exchange halves
- Repeat r times

DES (Tripple DES)

- Single DES breakable in less than 24h (complete search of key space)
- Tripple DES is still secure
- Three steps of DES on each data block using up to three keys
- Decryption in reverse sequence
- 3 independend keys are the most secure
- Three same keys can be used for (insecure) DES compatibility

AES Key Features

- FIPS standard 197
- Key length: 128/192/256 bit
- Block size: 128 bit
- Iterative rounds of substitutions and permutation, but no Feistel structure
- 10, 12 or 14 rounds
- Blocks of 16 bytes arranged in 4x4 state matrix
- Components of the round function are invertible and independent of key
 - **Substitute Bytes:** Non-linear substitution of bytes in state
 - **Shift Rows:** Cyclic shifting of rows
 - **Min Columns:** Multiplication of state elements with a fixed 4x4 matrix M

Modes of Operation for Block Cyphers

- Objective: Encrypt multiple plaintext blocks with the same block cypher
- Straightforward solution: blockwise encryption (“Electronic Codebook Mode”)
- Problem: Patterns in the distribution of plaintext blocks remain visible

Cypher Block Chaining (CBC)

- Avoids telltale patterns in ciphertext
- Decryption fails if a data block is missing or corrupted
- Each data block is encrypted in relation to the previous block

Counter Mode (CTR)

- Simple and efficient
- Random access still possible
- No issues if data block is missing
- Incrementing counter is involved in randomization per data block

Padding

- Plaintext needs to be a full number of blocks
- If plaintext does not fill the last block completely, it must be padded before encryption
 - In order to facilitate safe decryption, the last block is always padded: For example for a block size of n bytes, there are $1 \dots n$ bytes added to the plaintext before encryption
 - Decryption can check last bytes and strip them off correspondingly
- Always need to pad with at least one byte!
- Common methods
 - Pad with bytes of the same value as the number of padding bytes (PKCS#5; i.e. if there are three bytes to be padded, add 0x03 0x03 0x03)
 - Pad with 0x80 followed by 0x00 bytes
 - Pad with zeroes except for the last byte that indicates the number of padding bytes
 - Pad with zeroes
 - Pad with space characters (0x20)

Key Length Considerations

- Cryptography is always a matter of complexity
 - With enough time and/or space, all schemes can theoretically be broken
 - “brute force” attacks
 - Example: 56bit keys DES can be broken in <24h since 1999
- Meanwhile
 - 128bit keys have to be replaced in the coming years
 - 192bit keys are secure in medium term
 - 256bit keys are hard to crack due to physical boundaries
- Quantum computers might be able to crack keys much more quickly
- Numbers refer to unbroken algorithms in symmetric cryptography
 - Broken algorithm is one where an n bit key can be determined trying out significantly less than 2^n keys

Message Authentication

Message Authentication Codes (MACs)

- Objectives
 - **Integrity protection:** Prevent unauthorized manipulation of data
 - **Message authentication:** Prevent unauthorized origination on behalf of others
- Idea: Compute a cryptographic checksum (MAC)
- Required Properties
 - Cannot be counterfeited; without having the sender's secret, it is too complex to ...
 - * Find another message matching the same MAC
 - * Construct a suitable MAC for another message
 - Even smallest changes to message cause a big change of the MAC

General Scenario

Alice:

1. $m = \text{"I love you. Alice"}$
2. Select secret key K
3. Compute $MAC_K(m)$

Bob:

1. Receives m'
2. Selects secret key K
3. Computes $MAC_K(m')$
4. Compares computed MAC with received MAC \rightarrow Matches!

Assertion: If computed MAC equals the MAC included in the received message, an owner of the key (Alice) really sent this message and it was not changed on the way.

Scenario with Modified Message

Alice: Same as in General Scenario

Mallory:

- $m = \text{"It's all over! Alice."}$

Bob

1. Receives m'
2. Selects secret key K
3. Computes $MAC_K(m')$
4. Compares computed MAC with received MAC \rightarrow Doesn't match!
5. Ignore m

MAC Computation

- Requirements
 - Shared key k between sender and receiver
 - Hash function to create a code that changes if the message has been altered
- Using **block cypher** f_k and **hash function** $hash$: $MAC(m) = f_k(hash(m))$
- Using a **key dependent cryptographic hash function** $hash(k, m)$: $MAC(m) = hash(k, m)$

Hash Function Requirements

- Weak **collision resistance**: For a given message and hash it is impossible/to complex to find another message such that the hashes match
- **One-way** property
 - Easy to compute hash
 - Impossible to find message from hash

Asymmetric Encryption

Public Key Cryptography

Alice:

1. Generates key pair (PK_{Alice}, SK_{Alice})
2. Published PK_{Alice} at Trent's
3. c received \rightarrow decrypts $m = D_{SK_{Alice}}(C)$

Trent:

- Stores public keys
- Provides public keys on request

Bob:

1. Wants to send m to Alice confidentially
2. Obtains PK_{Alice} from Trent
3. Computes $c = E_{PK_{Alice}}(m)$
4. Sends c to Alice

RSA Key Generation

1. Alice chooses 2 large prime numbers p, q and computes $n = p \cdot q$, $\phi(n) = (p-1)(q-1)$
2. Alice chooses an integer e with $1 < e < \phi(n)$ that is relatively prime to $\phi(n)$
3. Alice computes an integer d with $1 < d < \phi(n)$ and $d \cdot e = k \cdot \phi(n) + 1$
4. Alice publishes her public key $PK_{Alice} = (e, n)$
5. Alice keeps her private key $SK_{Alice} = d$ and $p, q, \phi(n)$ secret

RSA Encryption

1. Bob obtains $PK_{Alice} = (e, n)$
2. Bob composes plaintext $m \in M = \{1, 2, \dots, n-1\}$
3. Bob computes the ciphertext $c = E_{PK_{Alice}}(m) = m^e \mod n$
4. Bob sends c to Alice

RSA Decryption

Alice can obtain the plaintext message m by computing $m = D_{SK_{Alice}}(c) = c^d \mod n = m^{ed} \mod n$

RSA Security

- **RSA problem:** Given e, n and $c = m^e \mod n$, find m
 - Most efficient approach to solve the RSA problem is currently the integer factorization of n : An upper limit to the complexity of the problem; can be used to derive the private key from the prime factors
 - Quantum computers will be more efficient in doing integer factorization (Shor's algorithm)
 - RSA problem and integer factorization still lack mathematical proof for their complexity
- **Organizational properties**
 - **Authenticity** of the public key (e, n)
 - **Confidentiality** of the secret key (d, p, q)
- **Mathematical properties**
 - **Complexity of factoring** the modulus n
 - **Complexity of solving** the RSA problem
- Failure of any properties will compromise the security of the method!

Hybrid Method

Combination of asymmetric and symmetric key methods.

Alice:

1. Generates key pair (PK_{Alice}, SK_{Alice})
2. Publishes PK_{Alice} at Trent's
3. c received \rightarrow Decrypts $K = D_{SK_{Alice}}(c)$
4. Alice and bob switch over to the symmetric key algorithm with key K

Trent:

- Stores public keys
- Provides public keys on request

Bob:

1. Obtains PK_{Alice} from Trent
2. Generates symmetric key K

3. Computes $c = E_{PK_{Alice}}(K)$
4. Sends c to Alice

Discrete Logarithms

Primitive element: Let p be a prime number. An element $g \leq p - 1$ is called primitive element $\mod p$ if for each $A \in \{1, 2, \dots, p - 1\}$ there is an x such that $A = g^x \mod p$

Discrete logarithm: Let p be a prime number and let $g \leq p - 1$ be a primitive element $\mod p$. Then an element x is called discrete logarithm of A to base $g \mod p$ if $A = g^x \mod p$.

Discrete logarithm problem: Given A, g, p , find $x \leq p - 1$ with $A = g^x \mod p$

One-Way Functions

- “Trap-door” functions
- Easy to compute in one direction (i.e. $f(x) = g^x \mod p$)
- Hard to invert
 - Ideally only possible using complete enumeration of all possible inputs
 - I.e. for a given y you need to try out all possible values $x = 0, 1, \dots, p - 1$ to find one $x_0 : f(x_0) = y$
- Definition of complexity often of the P and NP complexity classes
 - **P:** Answer of a problem can be found in polynomial time (b bits of problem size \rightarrow algorithm takes time b^k)
 - **NP:** Answer of problem cannot be found in polynomial time (b bits of problem size \rightarrow algorithm takes time k^b), but the correctness of given answer can be checked in polynomial time

Diffie-Hellman Key Exchange Protocol

Purpose: Allow communication partners without prior knowledge of another to establish a shared secret key over an insecure communication channel

1. Alice and Bob agree publicly on prime number p and a primitive element $g \leq p - 1$
2. Alice randomly chooses $\alpha \in \{2, \dots, p - 2\}$ and computes $A = g^\alpha \mod p$
3. Bob randomly chooses $\beta \in \{2, \dots, p - 2\}$ and computes $B = g^\beta \mod p$
4. Alice and Bob publicly exchange A and B
5. Alice and Bob hold a common secret key K :
 1. $K_B = A^\beta \mod p = g^{\alpha\beta} \mod p$
 2. $K_A = B^\alpha \mod p = g^{\alpha\beta} \mod p = K_B$

Diffie-Hellman Key Exchange Protocol Security

It depends on three properties which can't be relaxed:

- **Discrete logarithm problem:** There is no efficient inversion for integer exponentiation
- **Authenticity** of exchanged messages: No protection against MITM attacks!
- **Diffie-Hellman problem complexity:** Given $g, p, A = g^x \mod p, B = g^y \mod p$ find $K = g^{xy} \mod p$

Digital Signatures

- Requirements
 - **Tamper-proof**
 - **Unambiguous attribution** of signature to signing person/identity
 - **Inseparable connection** between signature and signed document
 - **Non-repudiability** of signature
- Typical approach
 - Encrypt hash of document with secret key
 - Signature can be verified using the public key

Alice:

1. Generates key pair (PK_{Alice}, SK_{Alice})
2. Publishes PK_{Alice} at Trent's
3. Computes $sig_{Alice}(m) = E_{SK_{Alice}}(hash(m))$

Trent:

- Stores public keys
- Provides public keys on request

Bob:

1. Obtains PK_{Alice} from Trent
2. Computes $hash(m_{received})$
3. Decrypts signature $D_{PK_{Alice}}(sig_{received})$
4. Compares $hash(m_{received})$ to the received signed hash

RSA Signatures

- Conventions
 - $PK_{Alice} = (e, n)$
 - $SK_{Alice} = d$
 - m is the message to be signed
 - h is the secure hash function
- **Computation** of signature: $sig_{Alice}(m) = (h(m))^d \mod n$
- **Verification** of signature
 - Bob receives (m', sig')
 - Bob computes $h(m')$ and $(sig')^e \mod n$
 - If both match, the signature is verified